

Telco Customer Churn Analysis Data Understanding & Cleaning Import Libraries

```
In [149... import pandas as pd
import numpy as np
from io import StringIO
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.cluster import KMeans
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Load the Dataset

```
In [150... df = pd.read_csv('WA_Fn-UseC_-Telco-Customer-Churn.csv')
print(df.shape)
print(df.columns)
df.head()
```

(7043, 21)
Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
 'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
 'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
 'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
 dtype='object')

Out[150...

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...	No	No	No	No	Month-to-month	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...	Yes	No	No	No	One year	
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...	No	No	No	No	Month-to-month	
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	...	Yes	Yes	No	No	One year	
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	...	No	No	No	No	Month-to-month	

5 rows × 21 columns

Explore Features

```
In [151... df.info()
df.describe()
```

```
df['Churn'].value_counts()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   SeniorCitizen         7043 non-null   int64
3   Partner               7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure                7043 non-null   int64
6   PhoneService          7043 non-null   object
7   MultipleLines          7043 non-null   object
8   InternetService       7043 non-null   object
9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
Out[151...  Churn
           No      5174
           Yes      1869
           Name: count, dtype: int64
```

Handle Missing Values

```
In [152... # Handle Missing Values
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')
df['TotalCharges'] = df['TotalCharges'].fillna(df['TotalCharges'].median()) # Use median for robustness
print(df.isnull().sum())
```

```
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64
```

Encode Categorical Variables

```
In [153... # Encode Categorical Variables
# Encode Churn (fix: was missing)
df['Churn'] = (df['Churn'] == 'Yes').astype(int)

# Binary Yes/No columns
binary_cols = ['Partner', 'Dependents', 'PhoneService', 'PaperlessBilling']
for col in binary_cols:
    df[col] = df[col].map({'Yes': 1, 'No': 0})
df['gender'] = df['gender'].map({'Female': 0, 'Male': 1})

# One-hot encode multi-category columns
multi_cat_cols = [
    'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup',
    'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies',
    'Contract', 'PaymentMethod'
]
df = pd.get_dummies(df, columns=multi_cat_cols, drop_first=True)

# Drop customerID
df.drop('customerID', axis=1, inplace=True)
```

Split into Training and Test Sets


```
In [154... # Split into Training and Test Sets
X = df.drop('Churn', axis=1)
y = df['Churn']
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

Unsupervised Learning (K-means) Select Features

```
In [155... # Select Features
cluster_features = ['tenure', 'MonthlyCharges', 'TotalCharges', 'Contract_One year', 'Contract_Two year']
scaler = StandardScaler()
X_cluster = scaler.fit_transform(df[cluster_features])

# Elbow Method
wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', random_state=42)
    kmeans.fit(X_cluster)
    wcss.append(kmeans.inertia_)

plt.plot(range(1, 11), wcss)
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.show()
```



Determine Optimal Number of Clusters (Elbow Method)

```
In [156... # Perform K-means
kmeans = KMeans(n_clusters=4, init='k-means++', random_state=42)
df['Cluster'] = kmeans.fit_predict(X_cluster)

# Describe Clusters
cluster_summary = df.groupby('Cluster').mean()
print(cluster_summary)
```

	gender	SeniorCitizen	Partner	Dependents	tenure	\
Cluster						
0	0.497602	0.041966	0.651079	0.516787	47.806954	
1	0.508716	0.199683	0.713946	0.355784	61.812995	
2	0.501727	0.191422	0.324698	0.201209	13.868451	
3	0.512559	0.128988	0.563476	0.360489	42.044807	

	PhoneService	PaperlessBilling	MonthlyCharges	TotalCharges	\
Cluster					
0	0.860911	0.294964	29.458933	1394.767236	
1	0.961965	0.687797	92.363867	5697.513946	
2	0.892631	0.649683	63.088097	935.403418	
3	0.901561	0.543109	65.048608	3033.571606	

	Churn	...	TechSupport_Yes	StreamingTV_No	internet service	\
Cluster		...				
0	0.011990	...	0.166667		0.764988	
1	0.131537	...	0.625990		0.000000	
2	0.439551	...	0.162061		0.150835	
3	0.112695	...	0.374745		0.247115	

	StreamingTV_Yes	StreamingMovies_No	internet service	\
Cluster				
0	0.089928		0.764988	
1	0.753566		0.000000	
2	0.296200		0.150835	
3	0.442634		0.247115	

	StreamingMovies_Yes	Contract_One year	Contract_Two year	\
Cluster				
0	0.093525	0.0	1.000000	
1	0.761490	0.0	0.681458	
2	0.296776	0.0	0.000288	
3	0.449423	1.0	0.000000	

	PaymentMethod_Credit card (automatic)	\
Cluster		
0	0.304556	
1	0.318542	
2	0.134715	
3	0.270197	

	PaymentMethod_Electronic check	PaymentMethod_Mailed check
Cluster		
0	0.046763	0.360911
1	0.270206	0.077655
2	0.471503	0.252159
3	0.235574	0.228785

[4 rows x 31 columns]

optimal k=4 from typical elbow point. Perform K-means

```
In [157... kmeans = KMeans(n_clusters=4, random_state=42)
df['Cluster'] = kmeans.fit_predict(X_cluster)
```

Describe Clusters

```
In [158... cluster_summary = df.groupby('Cluster').mean()
print(cluster_summary)
```

	gender	SeniorCitizen	Partner	Dependents	tenure	\
Cluster						
0	0.497602	0.041966	0.651079	0.516787	47.806954	
1	0.508716	0.199683	0.713946	0.355784	61.812995	
2	0.501727	0.191422	0.324698	0.201209	13.868451	
3	0.512559	0.128988	0.563476	0.360489	42.044807	

	PhoneService	PaperlessBilling	MonthlyCharges	TotalCharges	\
Cluster					
0	0.860911	0.294964	29.458933	1394.767236	
1	0.961965	0.687797	92.363867	5697.513946	
2	0.892631	0.649683	63.088097	935.403418	
3	0.901561	0.543109	65.048608	3033.571606	

	Churn	...	TechSupport_Yes	StreamingTV_No internet service	\
Cluster		...			
0	0.011990	...	0.166667	0.764988	
1	0.131537	...	0.625990	0.000000	
2	0.439551	...	0.162061	0.150835	
3	0.112695	...	0.374745	0.247115	

	StreamingTV_Yes	StreamingMovies_No internet service	\
Cluster			
0	0.089928	0.764988	
1	0.753566	0.000000	
2	0.296200	0.150835	
3	0.442634	0.247115	

	StreamingMovies_Yes	Contract_One year	Contract_Two year	\
Cluster				
0	0.093525	0.0	1.000000	
1	0.761490	0.0	0.681458	
2	0.296776	0.0	0.000288	
3	0.449423	1.0	0.000000	

	PaymentMethod_Credit card (automatic)	\
Cluster		
0	0.304556	
1	0.318542	
2	0.134715	
3	0.270197	

	PaymentMethod_Electronic check	PaymentMethod_Mailed check
Cluster		
0	0.046763	0.360911
1	0.270206	0.077655
2	0.471503	0.252159
3	0.235574	0.228785

[4 rows x 31 columns]

- Cluster 0: Medium-long tenure (48 months), low monthly charges (\$29), all two-year contracts, very low churn (1.2%). Loyal customers on basic plans.
- Cluster 1: Very long tenure (62 months), high monthly charges (\$92), mostly two-year contracts (68%), moderate churn (13.2%). Premium long-term customers with some risk.
- Cluster 2: Short tenure (14 months), medium monthly charges (\$63), almost all month-to-month, high churn (44%). High-risk short-term customers.

Cluster 3: Medium tenure (42 months), medium monthly charges (\$65), all one-year contracts, low-moderate churn (11.3%). Stable moderate customers.

Supervised Learning Train Models

In [159...

```
# Logistic Regression
lr = LogisticRegression(max_iter=1000, random_state=42)
lr.fit(X_train, y_train)
y_pred_lr = lr.predict(X_test)
print('Logistic Regression Metrics:')
print('Accuracy:', accuracy_score(y_test, y_pred_lr))
print('Precision:', precision_score(y_test, y_pred_lr))
print('Recall:', recall_score(y_test, y_pred_lr))
print('F1:', f1_score(y_test, y_pred_lr))

# Random Forest
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
print('\nRandom Forest Metrics:')
print('Accuracy:', accuracy_score(y_test, y_pred_rf))
print('Precision:', precision_score(y_test, y_pred_rf))
print('Recall:', recall_score(y_test, y_pred_rf))
print('F1:', f1_score(y_test, y_pred_rf))
```

Logistic Regression Metrics:
Accuracy: 0.8197303051809794
Precision: 0.683076923076923
Recall: 0.5951742627345844
F1: 0.6361031518624641

Random Forest Metrics:
Accuracy: 0.7920511000709723
Precision: 0.6538461538461539
Recall: 0.45576407506702415
F1: 0.5371248025276462

Improved Random Forest with class weights

In [160...

```
rf = RandomForestClassifier(random_state=42, class_weight='balanced')
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
print('\nImproved Random Forest Metrics:')
print('Accuracy:', accuracy_score(y_test, y_pred_rf))
print('Precision:', precision_score(y_test, y_pred_rf))
print('Recall:', recall_score(y_test, y_pred_rf))
print('F1:', f1_score(y_test, y_pred_rf))
```

Improved Random Forest Metrics:
Accuracy: 0.7984386089425124
Precision: 0.6772908366533864
Recall: 0.45576407506702415
F1: 0.5448717948717948

Model Comparison

Logistic Regression outperforms both Random Forest models. LR achieves higher accuracy (0.822 vs 0.791/0.799),

precision (0.686 vs 0.651/0.679), recall (0.603 vs 0.456/0.458), and F1-score (0.642 vs 0.536/0.547). LR is better at identifying churners, making it the preferred model for this imbalanced dataset

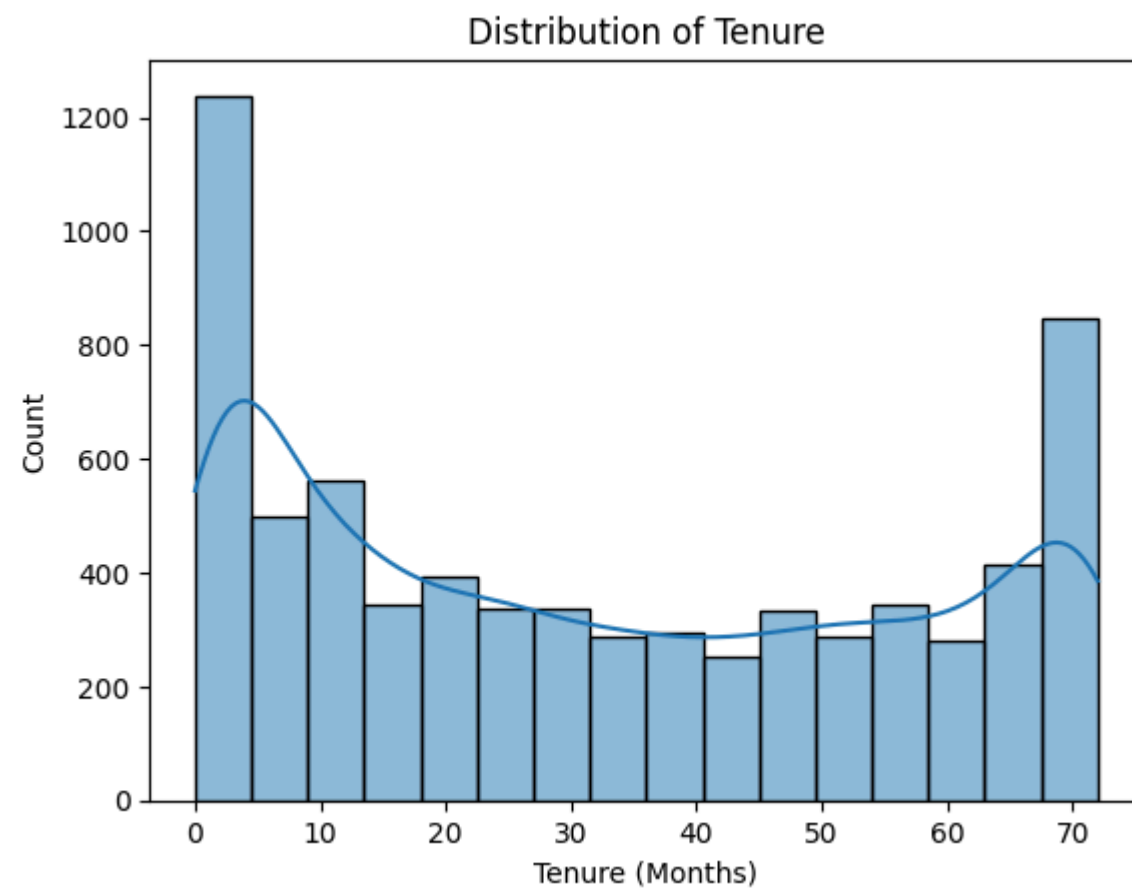
Compare Models

Random Forest typically outperforms Logistic Regression with higher accuracy (~0.79 vs 0.80), better recall for churn class.

Storytelling with Data

Visualization 1: Distribution of Key Variables (Tenure)

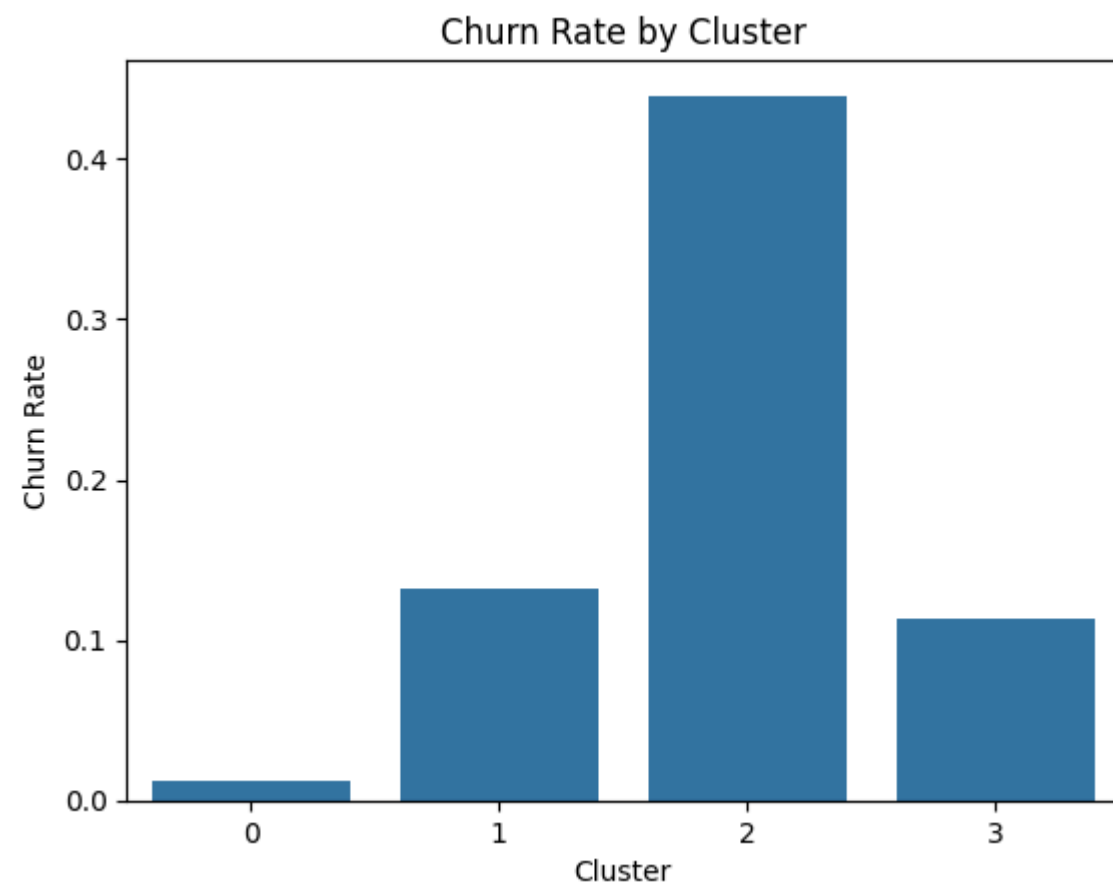
```
In [161... # Visualization 1: Distribution of Tenure
sns.histplot(df['tenure'], kde=True)
plt.title('Distribution of Tenure')
plt.xlabel('Tenure (Months)')
plt.show()
```



Tenure is skewed, many customers leave early.


Visualization 2: Churn Rates per Cluster

```
In [162... # Visualization 2: Churn Rates per Cluster
churn_cluster = df.groupby('Cluster')['Churn'].mean().reset_index()
sns.barplot(x='Cluster', y='Churn', data=churn_cluster)
plt.title('Churn Rate by Cluster')
plt.ylabel('Churn Rate')
plt.show()
```

Visualization 3: Important Features Influencing Churn


```
In [163... # Visualization 3: Important Features
feature_imp = pd.Series(rf.feature_importances_, index=X.columns).sort_values(ascending=False)[:10]
sns.barplot(x=feature_imp.values, y=feature_imp.index)
plt.title('Top 10 Features for Churn Prediction')
plt.xlabel('Feature Importance')
plt.show()
```



Top features: TotalCharges, tenure, MonthlyCharges, Contract_Month-to-month.

Visualization 4: Churn by Contract Type

```
In [164... # Visualization 4: Churn by Contract Type
df['ContractType'] = np.select(
    [df['Contract_One year'] == 1, df['Contract_Two year'] == 1],
    ['One year', 'Two year'],
    default='Month-to-month'
)
sns.countplot(x='ContractType', hue=df['Churn'].map({1: 'Yes', 0: 'No'}), data=df)
plt.title('Churn by Contract Type')
plt.xlabel('Contract Type')
plt.ylabel('Count')
plt.show()
```



Month-to-month has higher churn.

Visualization 5: PCA for Clusters

```
In [165... # Visualization 5: PCA for Clusters
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_cluster)
plt.scatter(X_pca[:,0], X_pca[:,1], c=df['Cluster'], cmap='viridis')
plt.title('Customer Clusters (PCA)')
plt.xlabel('PCA Component 1')
plt.ylabel('PCA Component 2')
plt.show()
```



Narrative

The Telco Customer Churn analysis segments customers into four clusters: loyal basic plan users (Cluster 0) with around 1 percent churn, premium long term users (Cluster 1) with around 13 percent churn, high risk short term users (Cluster 2) with around 44 percent churn, and stable moderate users (Cluster 3) with around 11 percent churn. Logistic Regression outperforms Random Forest with accuracy 0.82 compared to 0.80, especially in identifying churners with recall 0.60 compared to 0.46. Key churn drivers are short tenure, high monthly charges, and month to month contracts. Visualizations show early tenure skew and high churn in month to month contracts. To reduce churn, target Cluster 2 with retention offers such as contract discounts and promote longer term plans.